



DEEP LEARNING IN PYTHON

Understanding model optimization

Why optimization is hard

- Simultaneously optimizing 1000s of parameters with complex relationships
- Updates may not improve model meaningfully
- Updates too small (if learning rate is low) or too large (if learning rate is high)

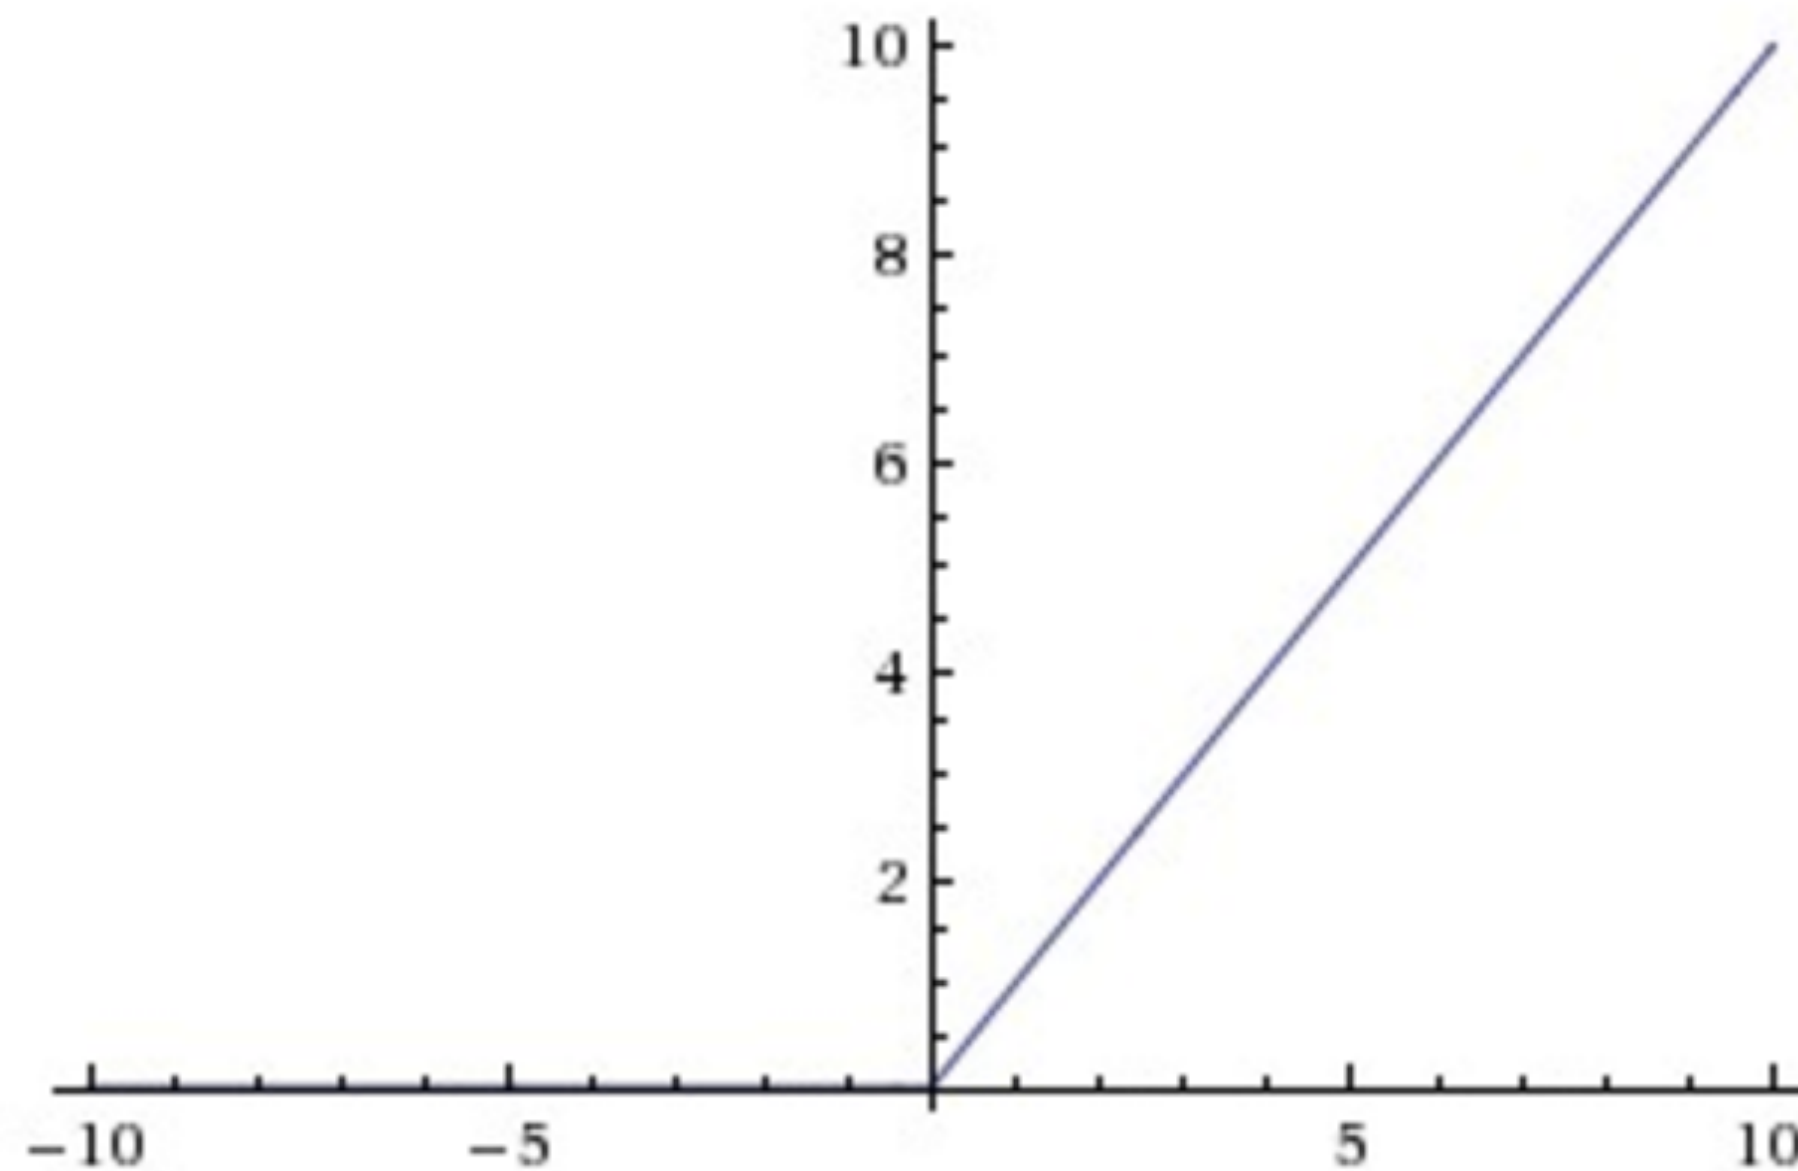



Stochastic gradient descent

```
In [1]: def get_new_model(input_shape = input_shape):  
  
In [2]:     model = Sequential()  
  
In [3]:     model.add(Dense(100, activation='relu', input_shape = input_shape))  
  
In [4]:     model.add(Dense(100, activation='relu'))  
  
In [5]:     model.add(Dense(2, activation='softmax'))  
  
In [6]:     return(model)  
  
In [7]: lr_to_test = [.000001, 0.01, 1]  
  
In [8]: for lr in lr_to_test:  
  
In [9]:     model = get_new_model()  
  
In[10]:     my_optimizer = SGD(lr=lr)  
  
In[11]:     model.compile(optimizer=my_optimizer, loss='categorical_crossentropy')  
  
In[12]:     model.fit(predictors, target)
```



The dying neuron problem




The dying neuron problem

- Once a node starts always getting negative inputs
 - It may continue only getting negative inputs
- Contributes nothing to the model
 - “Dead” neuron



Vanishing gradients



tanh function

Vanishing gradients

- Occurs when many layers have very small slopes (e.g. due to being on flat part of tanh curve)
- In deep networks, updates to backprop were close to 0



DEEP LEARNING IN PYTHON

Let's practice!



DEEP LEARNING IN PYTHON

Model validation

Validation in deep learning

- Commonly use validation split rather than cross-validation
- Deep learning widely used on large datasets
- Single validation score is based on large amount of data, and is reliable
- Repeated training from cross-validation would take long time



Model validation

```
In [1]: model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics=['accuracy'])
```

```
In [2]: model.fit(predictors, target, validation_split=0.3)
```

```
Epoch 1/10
```

```
89648/89648 [=====] - 3s - loss: 0.7552 - acc: 0.5775 - val_loss: 0.6969 -  
val_acc: 0.5561
```

```
Epoch 2/10
```

```
89648/89648 [=====] - 4s - loss: 0.6670 - acc: 0.6004 - val_loss: 0.6580 -  
val_acc: 0.6102
```

```
...
```

```
Epoch 8/10
```

```
89648/89648 [=====] - 5s - loss: 0.6578 - acc: 0.6125 - val_loss: 0.6594 -  
val_acc: 0.6037
```

```
Epoch 9/10
```

```
89648/89648 [=====] - 5s - loss: 0.6564 - acc: 0.6147 - val_loss: 0.6568 -  
val_acc: 0.6110
```

```
Epoch 10/10
```

```
89648/89648 [=====] - 5s - loss: 0.6555 - acc: 0.6158 - val_loss: 0.6557 -  
val_acc: 0.6126
```



Early Stopping

```
In [3]: from keras.callbacks import EarlyStopping  
  
In [4]: early_stopping_monitor = EarlyStopping(patience=2)  
  
In [5]: model.fit(predictors, target, validation_split=0.3, epochs=20,  
    ....:           callbacks = [early_stopping_monitor])
```



Output from early stopping

```
Train on 89648 samples, validate on 38421 samples
```

```
Epoch 1/20
```

```
89648/89648 [=====] - 5s - loss: 0.6550 - acc: 0.6151 - val_loss: 0.6548 -  
val_acc: 0.6151
```

```
Epoch 2/20
```

```
89648/89648 [=====] - 6s - loss: 0.6541 - acc: 0.6165 - val_loss: 0.6537 -  
val_acc: 0.6154
```

```
...
```

```
Epoch 8/20
```

```
89648/89648 [=====] - 6s - loss: 0.6527 - acc: 0.6181 - val_loss: 0.6531 -  
val_acc: 0.6160
```

```
Epoch 9/20
```

```
89648/89648 [=====] - 7s - loss: 0.6524 - acc: 0.6176 - val_loss: 0.6513 -  
val_acc: 0.6172
```

```
Epoch 10/20
```

```
89648/89648 [=====] - 6s - loss: 0.6527 - acc: 0.6176 - val_loss: 0.6549 -  
val_acc: 0.6134
```

```
Epoch 11/20
```

```
89648/89648 [=====] - 6s - loss: 0.6522 - acc: 0.6178 - val_loss: 0.6517 -  
val_acc: 0.6169
```

Experimentation

- Experiment with different architectures
- More layers
- Fewer layers
- Layers with more nodes
- Layers with fewer nodes
- Creating a great model requires experimentation



DEEP LEARNING IN PYTHON

Let's practice!




DEEP LEARNING IN PYTHON

Thinking about model capacity



Overfitting and underfitting



Workflow for optimizing model capacity

- Start with a small network
- Get the validation score
- Keep increasing capacity until validation score is no longer improving



Sequential experiments

Hidden Layers	Nodes Per Layer	Mean Squared Error	Next Step
1	100	5.4	Increase Capacity
1	250	4.8	Increase Capacity
2	250	4.4	Increase Capacity
3	250	4.5	Decrease Capacity
3	200	4.3	Done



DEEP LEARNING IN PYTHON

Let's practice!




DEEP LEARNING IN PYTHON

Stepping up to images



Recognizing handwritten digits

- MNIST dataset
- 28 x 28 grid flattened to 784 values for each image
- Value in each part of array denotes darkness of that pixel





DEEP LEARNING IN PYTHON

Let's practice!



DEEP LEARNING IN PYTHON

Final thoughts

Next steps

- Start with standard prediction problems on tables of numbers
- Images (with convolutional neural networks) are common next steps
- keras.io for excellent documentation
- Graphical processing unit (GPU) provides dramatic speedups in model training times
- Need a CUDA compatible GPU
- For training on using GPUs in the cloud look here:
[Link coming soon]



DEEP LEARNING IN PYTHON

Congratulations!